

Reinforcement Learning Tutorial


Instructor: Kanan Rahimli

In this project, I trained the Reinforcement Learning Agent to find its way to save the princess on a 5x5 maze. The code is simple and self explanatory but I will point out important parts. The code is written in a way that it can be used for any sized maze regardless of the position of elements:

Size = 25

New paths can be added:

```
allowed_paths = ((0,5), (5,6), (6,11), (10,11), (10,15), (15,20), (20,21),  
                (16,21), (16,17), (17,22), (22,23), (18,23), (13,18), (12,13),  
                (7,12), (7,8), (1,2), (2,3), (3,8), (3,4), (4,9), (9,14), (14,19),  
                (19,24))
```



For example, on the map there is a wall from 18 to 19 but we can “break” the wall by appending (18, 19) to the list.

We create a reward table filling in with dummy numbers, -1 in this case.

```
R = np.matrix(np.ones([size,size]))
```

```
R *= -1
```

The agent can 'decide to stay at its current position:

for i in range(size):

$R[i,i] = 0$

and the cost is 0

going back is allowed too:

for i in allowed_paths:

$a = i[:-1]$

 new_paths.append(a)

allowed_paths = allowed_paths + new_paths

and they cost 0 as well

bumping into bad guy costs -100

then we create a Q matrix (memory of the agent) filling with 0, since agent doesn't know anything at first

given a state, we return all the available actions from that state:

def available_actions(state):

 current_state_row = R[state,]

 av_act = np.where(current_state_row != -1)[1]

 return av_act

given all the possible actions, we choose one at random:

def sample_next_action(available_actions_range):

 next_action = int(np.random.choice(available_actions_range,1))

 return next_action

using the bellman formula, we update the q table, $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$:

def update(current_state, action, gamma):

 max_index = np.where(Q[action,] == np.max(Q[action,]))[1]

 if max_index.shape[0] > 1:

 max_index = int(np.random.choice(max_index, size = 1))

else:

 max_index = int(max_index)

max_value = Q[action, max_index]

Q[current_state, action] = R[current_state, action] + gamma * max_value

train the agent many times:

for i in range(size*1000):

 current_state = np.random.randint(0, int(Q.shape[0]))

 available_act = available_actions(current_state)

 action = sample_next_action(available_act)

 update(current_state, action, gamma)

we run the update function until we find princess, as in we give the agent his meaning of life:

while current_state != princess:

 do all the steps...

```
..      ..      ..      ..      ..      ..  
The best path to save the princess is:  
[11, 10, 15, 20, 21, 16, 17, 22, 23, 18, 13, 12, 7, 8, 3, 4, 9, 14, 19, 24]  
>>>
```